

CS 540

Database Management Systems

Relational Schema Design

Integrity constraints

- We like to restrict the data stored in the database.
 - ssn contains only numerical symbols.
 - name does **not** contain numerical symbols.
 - Tuples with equal ssn values have equal names.

Emp

ssn	name	address
111122222	John	21 Kings St.
111122222	John	234 2 nd St.
454565761	Charles	31 Kings St.
454565761	Charles	2 Harrison St.

- We use **integrity constraints** to express these restrictions.

Functional dependency (FD) & key

- Given set of attributes X, Y in relation R, the **functional dependency** $X \rightarrow Y$ means that all tuples in R that agree on attributes in X must also agree on Y.

Emp

ssn \rightarrow name? Yes

ssn \rightarrow address ? No

ssn	name	address
111122222	John	21 Kings St.
111122222	John	234 2 nd St.
454565761	Charles	31 Kings St.
454565761	Charles	2 Harrison St.

- A **key** in R is a set of attributes of R that functionally determines *all attributes* in R and *none of its subsets* is a key.
 - ssn? {ssn, address}? {ssn, name, address}?
- Super-key** is a set of attributes that contains a key.

How to find FDs and keys?

- From domain and domain experts.
- Logically implied by other FDs.
- *Example:*
`movies(title, year, actor, cost, revenue, b-buster)`
 $FD_1: title, year, actor \rightarrow cost$
 $FD_2: title, year, actor \rightarrow revenue$
Implied $FD_3: title, year, actor \rightarrow cost, revenue$
- **Closure of a set of FDs:** all FDs implied by a set of FDs.
- Generate them using Armstrong's axioms.

Armstrong's axioms

- **Reflexivity:** $A_1, \dots, A_n \rightarrow A_1$
generally, $A_1, \dots, A_n \rightarrow A_i, \dots, A_j$; $1 \leq i, j \leq n$ (*trivial FD*)
- **Augmentation:**
If $A_1, \dots, A_n \rightarrow B_1, \dots, B_m$ then
 $A_1, \dots, A_n, C_1, \dots, C_k \rightarrow B_1, \dots, B_m, C_1, \dots, C_k$
- **Transitivity:**
If $A_1, \dots, A_n \rightarrow B_1, \dots, B_m$ and $B_1, \dots, B_m \rightarrow C_1, \dots, C_k$
then $A_1, \dots, A_n \rightarrow C_1, \dots, C_k$

Computing the closure of a set of FDs (U)

- $U^+ = U$.
- Repeat
 - Apply **reflexivity** and **augmentation** to each FD in U^+ and add the resulting FDs to U^+ .
 - Apply **transitivity** to each pairs of FDs in U^+ and add the resulting FDs to U^+ .
- Until U^+ does not change anymore.

Useful rules

- They are derived from axioms and we may use them to derive implied FDs faster.

- **Decomposition**

If $A_1, \dots, A_n \twoheadrightarrow B_1, \dots, B_m$ then $A_1, \dots, A_n \twoheadrightarrow B_1$,
 $A_1, \dots, A_n \twoheadrightarrow B_2, \dots$, and $A_1, \dots, A_n \twoheadrightarrow B_m$.

Proof: ?

- **Union**

If $A_1, \dots, A_n \twoheadrightarrow B_1, \dots$, and $A_1, \dots, A_n \twoheadrightarrow B_m$ then
 $A_1, \dots, A_n \twoheadrightarrow B_1, \dots, B_m$.

Proof: ?

Computing the closure of a set of FDs

`movies(title, year, actor, cost, revenue, b-buster)`

FD₁ title, year, actor → cost

FD₂ title, year, actor → revenue

FD₃ cost, revenue → b-buster

Apply union on FD₁ and FD₂:

FD₄: title, year, actor → cost, revenue

Apply transitivity on FD₄ and FD₃:

FD₅: title, year, actor → b-buster

Apply union on FD₄ and FD₅:

FD₅: title, year, actor → cost, revenue, b-buster

Enforcing FDs

- Update '*John*' to '*Richard*' in the first tuple.
 - violates $ssn \rightarrow name$.
 - update anomaly.

Emp

ssn	name	address
111122222	John	21 Kings St.
111122222	John	234 2 nd St.
454565761	Charles	31 Kings St.
454565761	Charles	2 Harrison St.

- Write a program that checks the database after each update for violations of all FDs.
 - Large relations and many FDs => inefficient.

More problems

- Delete John's addresses => **lose his ssn and name.**
 - **deletion anomaly.**
- Insert a tuple with new ssn and name **but no address.**
 - **insertion anomaly.**

Emp

ssn	name	address
111122222	John	21 Kings St.
111122222	John	234 2 nd St.
454565761	Charles	31 Kings St.
454565761	Charles	2 Harrison St.

- One may use NULL values to solve these problems.
 - They do not necessarily indicate lack of knowledge.
 - It is hard to write correct SQL queries over the database.

Normalization

- Transform the schema to a schema without update, deletion, and insertion anomalies.

Emp

ssn	name	address
111122222	John	21 Kings St.
111122222	John	234 2 nd St.
454565761	Charles	31 Kings St.
454565761	Charles	2 Harrison St.



Emp-name

ssn	name
111122222	John
454565761	Charles

Emp-addr

ssn	address
111122222	21 Kings St.
111122222	234 2 nd St.
454565761	31 Kings St.
454565761	2 Harrison St.

- update anomaly?
- deletion anomaly?
- insertion anomaly?

Normalization

- Given schema S_1 , find schema S_2 that does not have update/deletion/insertion anomalies.

Emp

ssn	name	address
111122222	John	21 Kings St.
111122222	John	234 2 nd St.
454565761	Charles	31 Kings St.
454565761	Charles	2 Harrison St.



Emp-name

ssn	name
111122222	John
454565761	Charles

Emp-addr

ssn	address
111122222	21 Kings St.
111122222	234 2 nd St.
454565761	31 Kings St.
454565761	2 Harrison St.

$S_1 = (\{\text{Emp}\}, \{\text{ssn} \rightarrow \text{name}\})$

$S_2 = (\{\text{Emp-name}, \text{Emp-addr}\}, \{\text{ssn} \rightarrow \text{name}\})$

- Schema S_2 is in a **normal form**.

Many normal forms

	Normal form	Defined by	In
1NF	First normal form	Two versions: E.F. Codd (1970), C.J. Date (2003)	1970 ^[1] and 2003 ^[9]
2NF	Second normal form	E.F. Codd	1971 ^[2]
3NF	Third normal form	Two versions: E.F. Codd (1971), C. Zaniolo (1982)	1971 ^[2] and 1982 ^[10]
EKNF	Elementary Key Normal Form	C. Zaniolo	1982 ^[10]
BCNF	Boyce–Codd normal form	Raymond F. Boyce and E.F. Codd	1974 ^[11]
4NF	Fourth normal form	Ronald Fagin	1977 ^[12]
5NF	Fifth normal form	Ronald Fagin	1979 ^[13]
DKNF	Domain/key normal form	Ronald Fagin	1981 ^[14]
6NF	Sixth normal form	C.J. Date, Hugh Darwen, and Nikos Lorentzos	2002 ^[15]

Boyce-Codd normal form (BCNF)

- Relation R is in **BCNF**, if and only if:
 - For each non-trivial FD $X \rightarrow Y$, X is a super-key of R.
- Every attribute depends only on super-keys.

- **Example:** $ssn \rightarrow name$
 - ssn is **not** a super-key.
 - *Employee* is **not** in BCNF.

Emp

ssn	name	address
111122222	John	21 Kings St.
111122222	John	234 2 nd St.
454565761	Charles	31 Kings St.
454565761	Charles	2 Harrison St.

- Decompose *Emp*.
 - super-key of *Emp-name*?
 - super-key of *Emp-addr*?
 - Schema is in BCNF.

Emp-name

ssn	name
111122222	John
454565761	Charles

Emp-addr

ssn	address
111122222	21 Kings St.
111122222	234 2 nd St.
454565761	31 Kings St.
454565761	2 Harrison St.

BCNF decomposition of R

- Find the closure of FDs in R.
- Find the keys for R.
- Pick an FD $A \twoheadrightarrow B$ that violates the BCNF condition in R.
 - select the largest possible B.
- Decompose relation R to relational R1 and R2.

R	Other attributes	A	B
----------	------------------	---	---

R1	Other attributes	A	R2	A	B
-----------	------------------	---	-----------	---	---

- Repeat until there is no BCNF violation left.

BCNF decomposition example

`Emp2(ssn, name, street, city, state, zip)`

$FD_1 \text{ ssn} \rightarrow \text{name}$

$FD_2 \text{ zip} \rightarrow \text{state}$

Key: {ssn, street, city, zip}

FD_1 violates BCNF.

`Emp2(ssn, name)`

`Emp-addr(ssn, street, city, state, zip)`

FD_2 violates BCNF.

`Emp2(ssn, name)`

`Emp-addr(ssn, street, city, zip)`

`Location(zip, state)`

The danger of normalization

- Losing information.
- We like to preserve the information stored in R.
 - Recover the tuples in R from R's BCNF decomposition:
lossless decomposition
 - Recover all FDs in R from its BCNF decomposition:
dependency preserving
- Check these condition after normalizing a relation.

Lossless decomposition

- If $(R1, R2)$ is a decomposition of R , then $\text{Join}(R1, R2) = R$ for all instances of $R, R1, R2$.
- **Example:** $R(A, B, C) \rightarrow R1(A, B), R2(A, C)$

A	B	C
a1	b1	c1
a1	b2	c2



A	B
a1	b1
a1	b2

A	C
a1	c1
a1	c2

A	C
a1	c1
a1	c2

A	B
a1	b1
a1	b2

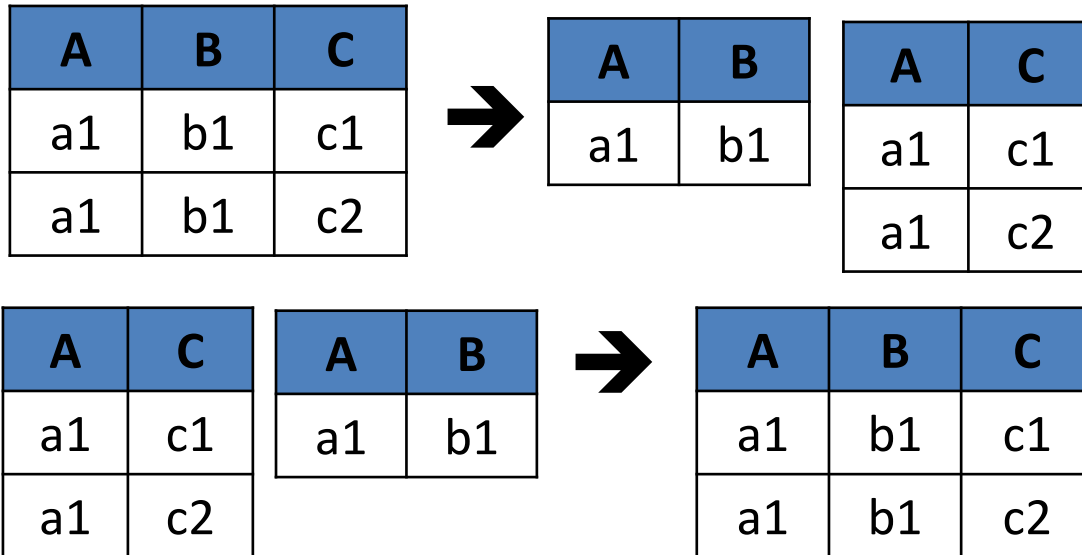


A	B	C
a1	b1	c1
a1	b1	c2
a1	b2	c1
a1	b2	c2

we get bogus tuples, not a lossless decomposition.

BCNF decomposition is lossless

- Example:** $\{R(A, B, C), A \twoheadrightarrow B\}$ $\{R_1(A, B), R_2(A, C), A \twoheadrightarrow B\}$



The join does not produce bogus tuples.

Informally speaking, if the join attribute is a key, we are OK.

It is true for all BCNF decompositions.

BCNF is **not** dependency preserving

$\text{Emp}(\text{ssn}, \text{name}, \text{address}), \text{ssn} \rightarrow \text{name}, \text{name}, \text{address} \rightarrow \text{ssn}$

BCNF decomposition:

$\text{Emp-name}(\text{ssn}, \text{name}) \text{ssn} \rightarrow \text{name}$

$\text{Emp-addr}(\text{ssn}, \text{address})$ No FD!

not dependency preserving

- What will happen?

Emp-name

ssn	name
111122222	John
444455555	John

Emp-addr

ssn	address
111122222	21 Kings St.
111122222	21 Kings St.
444455555	21 Kings St.
444455555	21 Kings St.



Emp

ssn	name	address
111122222	John	21 Kings St.
111122222	John	21 Kings St.
444455555	John	21 Kings St.
444455555	John	21 Kings St.

satisfies the dependencies

does **not** satisfies the dependencies

Given FD $A \rightarrow B$, If normalization puts A and B in different relations, it is not dependency preserving.

Dependency preserving normal form: 3NF

- Relation R is in 3rd normal form if for each non-trivial FD $X \rightarrow Y$ in R
 - X is a super-key or
 - Y is a part of a key.
- Every non-key attribute must depend on a key, the whole key, and nothing but the key!
- **Example:**
`Emp(ssn, name, address)`
 $ssn \rightarrow name, address, name \rightarrow ssn$
 - 3NF but not BCNF.
 - 3NF is a lossless decomposition and preserves dependencies.

Minimal basis (minimal cover)

- The FD sets **U1** and **U2** are *equivalent* if and only if $U1^+ = U2^+$.
- **U2** is a *minimal basis* for **U1** iff:
 - **U2** and **U1** are equivalent.
 - The FDs in **U2** have only one attribute in their right hand side.
 - If we remove an FD from **U2** or an attribute from an FD in **U2**, **U2** will not be equivalent to **U1**.
- **Intuition:** a smaller set of FDs with fewer attributes that implies U^+ .

Finding minimal basis of U

1. **Standard form:** replace each FD $A_1, \dots, A_n \rightarrow B_1, \dots, B_m$ by $A_1, \dots, A_n \rightarrow B_1, \dots, A_1, \dots, A_n \rightarrow B_m$
 2. **Minimize left hand side:** for each attribute A_j and each FD $A_1, \dots, A_j, \dots, A_n \rightarrow B_i$, check if you can remove A_j from the FD while preserving U^+ .
1. **Delete redundant FD:** check each remaining FD to see if you can remove it while preserving U^+ .
- The minimal basis for U is not necessarily unique.

Finding minimal basis of U

- **Example:**

$$U = \{B \rightarrow A, D \rightarrow A, AB \rightarrow D\}$$

1. It is already in standard form.
2. We can remove A from $\{AB \rightarrow D\}$ to get
$$U = \{B \rightarrow A, D \rightarrow A, B \rightarrow D\}$$
3. We can remove $B \rightarrow A$ as it is implied by transitivity from others.

3NF synthesizing algorithm

Input: relation R and set of FDs U.

Output: Normalized schema S

1. Find a minimal basis **M** for U.
 2. For each FD $A \rightarrow B$ in **M**, if AB is not covered by any relation in S, add $R_i = (A, B)$ to S.
 3. If none of the relations in S contain a super-key for R, add a relation to S whose attributes form a key for R.
- **Why step 3?** consider relation $R(A, B, C)$ with $U = \{A \rightarrow B, C \rightarrow B\}$.
Minimal basis of U is $U : 3NF = R_1(\underline{A}, B) R_2(\underline{C}, B)$
lossless join property? $3NF = R_1(\underline{A}, B) R_2(\underline{C}, B) R_3(\underline{A}, \underline{C})$

3NF versus BCNF

- BCNF eliminates more redundancies than 3NF.
- Normalization must be dependency preserving, unless there is a strong reason.
- Try BCNF, but if it is not dependency preserving use 3NF.

De-normalization

- Normalization improves data quality, but it has some drawbacks.
 - **performance**: normalized schemas require more joins.
 - **readability**: normalized schemas are hard to understand.
 - they contains larger numbers of relations.
 - they place related attributes in different relations.
- DB designers find the trade-off.
 - No write in OnLine Analytical Processing (OLAP) ➔ argues against normalization.
 - write in OnLine Transaction Processing (OLTP) ➔ argues for normalization.