

CS 540

Database Management Systems

Review of Relational Model and SQL

Review problems:

people betting on OSU football games

Out(game, outcome) Bets(who, outcome, game, amt)

game	outcome
USC	W
UCLA	L
Stanford	W

who	outcome	game	amt
John	W	USC	200
John	W	UCLA	100
John	L	Arizona	150
Kevin	L	UO	210
Kevin	L	UCLA	50
Kevin	W	Stanford	120

Some games have not been played yet, e.g., *Arizona*.

Problem 1

List the completed games that nobody bet on.

game	outcome
USC	W
UCLA	L
Stanford	W

game

who	outcome	game	amt
John	W	USC	200
John	W	UCLA	100
John	L	Arizona	150
Kevin	L	UO	210
Kevin	L	UCLA	50
Kevin	W	Stanford	120

Problem 1

List the completed games that nobody bet on.

(**Select** Game

From Out)

Except

(**Select** Game

From Bets)

Problem 2

Who bet the most money on a single game?

game	outcome
USC	W
UCLA	L
Stanford	W

who	amt
Kevin	210

who	outcome	game	amt
John	W	USC	200
John	W	UCLA	100
John	L	Arizona	150
Kevin	L	UO	210
Kevin	L	UCLA	50
Kevin	W	Stanford	120

Problem 2

Who bet the most money on a single game?

Select Who, Amt

From Bets

Where Amt >= **All**

(**Select** Amt

From Bets)

Problem 3

List the games that all bettors agree on.

game	outcome
USC	W
UCLA	L
Stanford	W

game
Stanford
Arizona
UO
USC

who	outcome	game	amt
John	W	USC	200
John	W	UCLA	100
John	L	Arizona	150
Kevin	L	UO	210
Kevin	L	UCLA	50
Kevin	W	Stanford	120

Problem 3

List the games that all bettors agree on.

(Select game

From Bets)

Except

(Select Bets1.game

From Bets Bets1, Bets Bets2

Where (Bets1.game = Bets2.game)

And (Bets1.outcome <> Bets2.outcome))

Problem 4

For each game, the number of people betting on OSU to win and the number betting on OSU to lose.

game	outcome
USC	W
UCLA	L
Stanford	W

game	outcome	num
Stanford	W	1
UO	L	1
UCLA	W	1
UCLA	L	1
USC	W	1
Arizona	L	1

who	outcome	game	amt
John	W	USC	200
John	W	UCLA	100
John	L	Arizona	150
Kevin	L	UO	210
Kevin	L	UCLA	50
Kevin	W	Stanford	120

Problem 4

For each game, the number of people betting on OSU to win and the number betting on OSU to lose.

```
Select game, outcome, Count(who) As num  
From Bets  
Group By game, outcome
```

Problem 5

Find the people who have made two or more bets on OSU to lose.

game	outcome
USC	W
UCLA	L
Stanford	W

who
Kevin

who	outcome	game	amt
John	W	USC	200
John	W	UCLA	100
John	L	Arizona	150
Kevin	L	UO	210
Kevin	L	UCLA	50
Kevin	W	Stanford	120

Problem 5

Find the people who have made two or more bets on OSU to lose.

Select who

From Bets

Where outcome = 'L'

Group By who

Having Count(outcome) >= 2

Problem 6

Who bet the most money overall?

game	outcome
USC	W
UCLA	L
Stanford	W

who	sumAmt
John	450

who	outcome	game	amt
John	W	USC	200
John	W	UCLA	100
John	L	Arizona	150
Kevin	L	UO	210
Kevin	L	UCLA	50
Kevin	W	Stanford	120

Problem 6

Who bet the most money overall?

```
Select who, Sum(amt) As sumAmt
From Bets
Group By who
Having Sum(amt) >= ALL
                                (Select Sum(amt)
                                From Bets
                                Group By who)
```

Problem 7

Who has bet on every game?

game	outcome
USC	W
UCLA	L
Stanford	W

who	outcome	game	amt
John	W	USC	200
John	W	UCLA	100
John	L	Arizona	150
Kevin	L	UO	210
Kevin	L	UCLA	50
Kevin	W	Stanford	120

Who

Problem 7

Who has bet on every game?

```
Create View AllGames As  
(Select game From Out) Union  
(Select game From Bets)
```

```
Select who  
From Bets  
Group By who  
Having Count(Distinct game) =  
(Select Count(Distinct game)  
  From AllGames)
```


Problem 8

What games have won the most money for the people who bet on OSU to win?

game	outcome
USC	W
UCLA	L
Stanford	W

game
USC

who	outcome	game	amt
John	W	USC	200
John	W	UCLA	100
John	L	Arizona	150
Kevin	L	UO	210
Kevin	L	UCLA	50
Kevin	W	Stanford	120

Problem 8

What games have won the most money for the people who bet on OSU to win?

Create View Success-Win **As**

(Select Bets.game, **Sum**(Bets.amt) **As** SumAmt

From Bets, Out

Where out.game = Bets.game **And**

Bets.outcome = 'W' **And** Out.outcome = 'W'

Group By Game)

Select Distinct game

From Success-Win

Where SumAmt >= **All**

(Select SumAmt

From Success-Win)

Problem 9

List the people who won some money so far.

game	outcome
USC	W
UCLA	L
Stanford	W

who
John
Kevin

who	outcome	game	amt
John	W	USC	200
John	W	UCLA	100
John	L	Arizona	150
Kevin	L	UO	210
Kevin	L	UCLA	50
Kevin	W	Stanford	120

Problem 9

List the people who won some money so far.

Create View Success As

```
(Select who, Sum(amt) As pAmt  
From Bets, Out  
Where Out.game = Bets.game And Bets.outcome = Out.outcome  
Group By who)
```

Create View Failure As

```
(Select who, Sum(amt) As nAmt  
From Bets, Out  
Where Out.game = Bets.game And Bets.outcome <> Out.outcome  
Group By who)
```

Problem 9

List the people who won some money so far.

```
(Select Success.who  
  From Success, Failure  
  Where Success.who = Failure.who And  
        Success.pAmt > Failure.nAmt )
```

Union

```
(Select who  
  From Success  
  Where who not in  
        (Select who  
          From Failure) )
```

Views

- Virtual relations defined over *base relations*

Create View <name> **As** <query>;

- PopularShop(sname, addr)*: shops with at least 10 frequent drinkers.

Create View PopularShop **As**

Select C.sname, C.addr

From CoffeeShop C, Frequents F

Where C.sname = F.sname

Group By C.sname, C.addr

Having Count(dname) > 9

- Used in queries like the base relations.

Select addr

From PopularShops;

Advantages of using views

- **Usability**
 - easier for users to find relevant relations.
 - DBAs create the relevant views for each group of users.
- **Security**
 - each group of users should access only a subset of database.
 - DBAs define views and their access privileges.
- **Efficiency**
 - One may precompute and store frequently used views.

Virtual versus materialized views

- **Virtual view**
 - database system stores the view definitions
 - database system rewrites the query using view definitions
- **Materialized view**
 - Database system pre-computes the views
 - Users issue many queries using *PopularShops*
 - **advantage:** improves the running time for some queries.
 - **disadvantage:** database system should refresh it if its base relations are modified.

Which views to materialize?

- **Tradeoff:** #read from view vs. #write to base relations
- Two types of database systems
 - Online Transaction Processing (OLTP)
 - banking systems, online shopping, ...
 - heavy write workload
 - Online Analytics Processing (OLAP)
 - integrated database about branches
 - data warehouses
 - write is not frequent, e.g., once a month
 - many complex join queries with aggregation functions
- Materialized views are popular in data warehouses
 - no significant overhead

Maintaining materialized views with both r/w

- No silver bullet!
 - choose a policy that matches the application needs.
- Two steps
 - **Propagate:** compute changes to the view when data changes.
 - **Refresh:** apply changes to the materialized view table.
- Maintenance policy
 - **Immediate:** as soon as the base tables are updated.
 - **advantage:** views are always consistent.
 - **disadvantage:** updates are slower.
 - **Deferred:** some time later
 - **advantage:** scale to many views without slowing down updates .
 - **disadvantage:** views may become inconsistent.

Deferred maintenance

- Lazy
 - Delay refresh until the next query over the view
 - Refresh before answering the query
- Periodic (snapshot)
 - refresh the view periodically
 - **disadvantage:** Some queries may be answered over an outdated version of view.
 - **advantage:** scale to update many views
 - more widely used

Support for view materialization

- Example in Oracle

```
Create materialized view PopularShop As  
  Select sname, addr  
  From CoffeeShop C, Frequents F  
  Where C.sname = F.sname  
  Group By C.sname, C.addr  
  Having Count(dname) > 9
```

- Also called *snapshot*.

- No direct support in MySQL

- programmers have to implement it

View selection

- Given a query workload, find the minimum set of views to materialize such that
 - minimizes the running time of all queries
- General approach
 - find the time-consuming joins /aggregation sub-queries.
 - pick the ones that are used in sufficiently many times.
- Some database systems provide wizard programs to help

Is SQL Sufficient?

- Using `IsParent(parent,child)`, find grand children of a given person.
- Now find all descendants of a given person.
- It is **not** possible to write this query in standard SQL!
 - We can prove it.
 - *3 bonus points for everybody who finds, understands and writes the main ideas of the proof and submits it by 2/1.*

Recursive queries in SQL

- SQL standards have recursive SQL
 - most database systems do not implement that
- Database systems usually support limited recursion
 - MySQL recursive cte, Oracle's connected by, ...
- They define within a single query
 - a base case (base query)
 - a recursion step
- Systems limit the type of queries used for recursion
 - not group by/ aggregation function
 - to keep the plan for normal queries fast.

Recursive queries in MySQL

- Common table expression (CTE)
 - relation variable within the scope of a single query

```
WITH cte (col1, col2) AS  
(  
    SELECT 1, 2  
    UNION ALL  
    SELECT 3, 4  
)  
SELECT col1, col2 FROM cte;
```


Recursive queries in MySQL

- Common table expression (CTE)

WITH

cte1 **AS** (**SELECT** a, b **FROM** table1),

cte2 **AS** (**SELECT** c, d **FROM** table2)

SELECT b, d

FROM cte1 **JOIN** cte2

WHERE cte1.a = cte2.c;

- used similar to virtual view in the query
 - the query plan may be more efficient as each CTE is executed only once and used multiple times.

Recursive queries in MySQL

- Recursive CTE

```
WITH RECURSIVE cte (n) AS
(
    SELECT 1
    UNION ALL
    SELECT n + 1 FROM cte WHERE n < 5
)
SELECT * FROM cte;
```

base case (base query) ?

recursion step ?

Recursive queries in MySQL

- Using *Employee(id, name, manager_id)* produce the organizational chart of the management chain.

```
WITH RECURSIVE employee_paths (id, name, path) AS
(
    SELECT id, name, CAST(id AS CHAR(200))
    FROM employees WHERE manager_id IS NULL
    UNION ALL
    SELECT e.id, e.name, CONCAT(ep.path, ',', e.id)
    FROM employee_paths AS ep
    JOIN employees AS e ON ep.id = e.manager_id
)
SELECT * FROM employee_paths ORDER BY path;
```