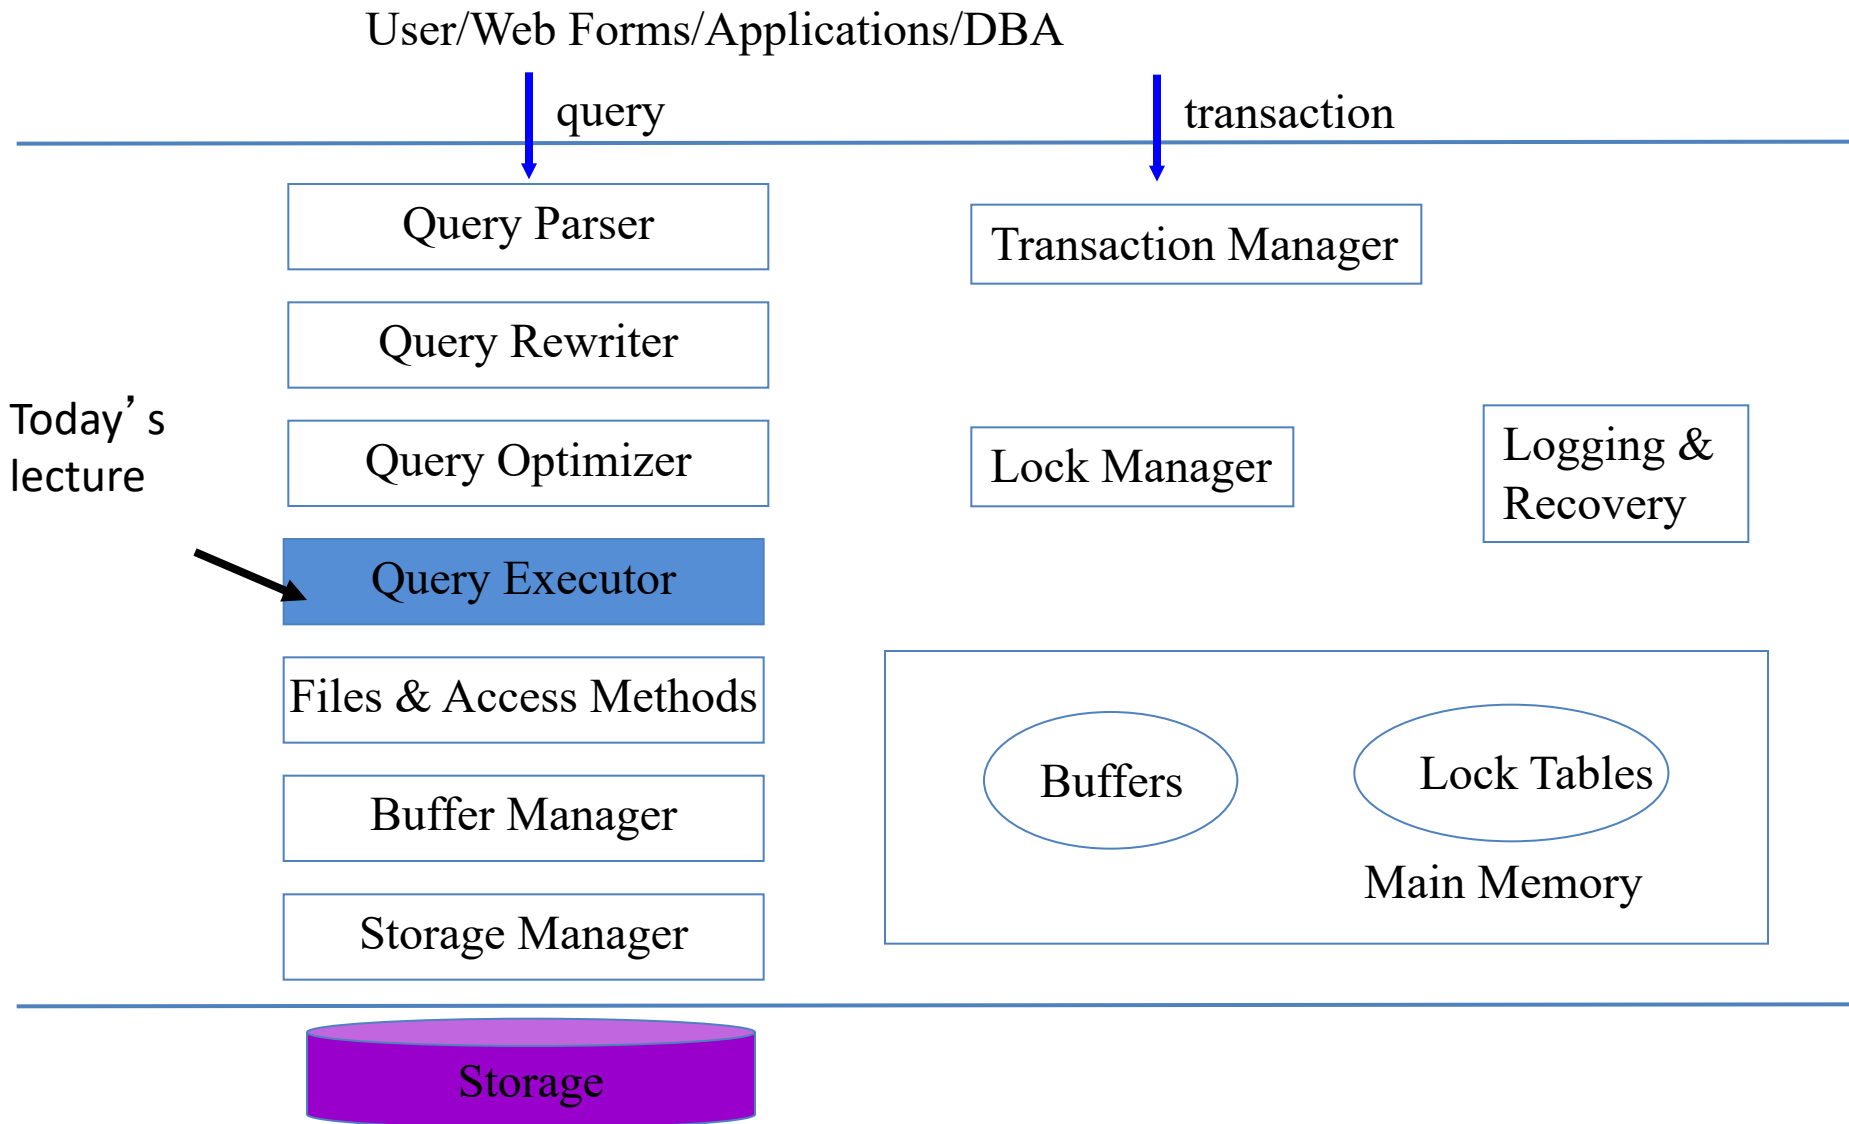


CS 540

Database Management Systems

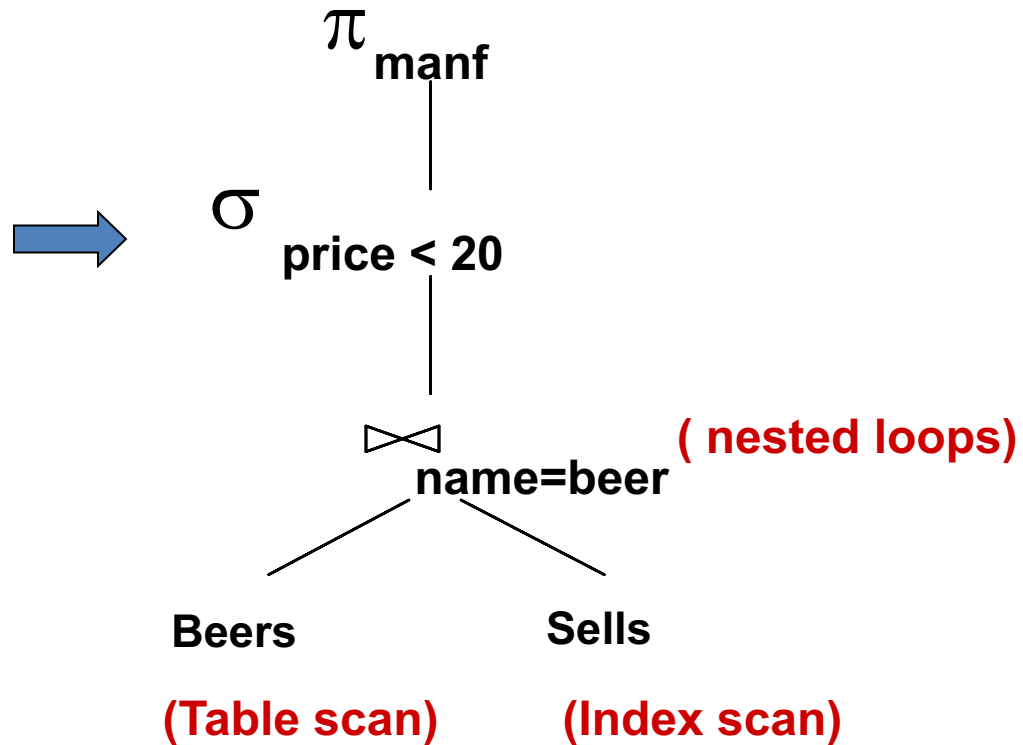
Query Processing

DBMS Architecture



Query Execution Plans

```
SELECT B.manf
FROM   Beers B, Sells S
WHERE  B.name=S.beer AND
       S.price < 20
```



Query Plan:


- logical plan (declarative)
- physical plan (procedural)
 - procedural implementation of each logical operator
 - scheduling of operations

Logical versus physical operators

- Logical operators
 - Relational Algebra Operators
 - Join, Selection, Projection, Union, ...
- Physical operators
 - Algorithms to implement logical operators.
 - Hash join, nested loop join, ...
- More than one physical operator for each logical operator

Explain command in Postgres

```
SELECT C.STATE, SUM(O.NETAMOUNT), SUM(O.TOTALAMOUNT)
FROM CUSTOMERS C
    JOIN CUST_HIST CH ON C.CUSTOMERID = CH.CUSTOMERID
    JOIN ORDERS O ON CH.ORDERID = O.ORDERID
GROUP BY C.STATE
```



Communication between operators: iterator model

- Each physical operator implements three functions:
 - **Open**: initializes the data structures.
 - **GetNext**: returns the next tuple in the result.
 - **Close**: ends the operation and frees the resources.
- It enables pipelining
- Other option: compute the result of the operator in full and store it in disk or memory:
 - inefficient.

Physical operators

- Logical operator: **selection**
 - read the entire or selected tuples of relation R.
 - tuples satisfy some predicate
- **Table-scan:** R resides in the secondary storage, read its blocks one by one.
- **Index-scan:** If there is an index on R, use the index to find the blocks.
 - more efficient
- Other operators for *join*, *union*, *group by*, ...
 - **join is the most important one.**
 - **focus of our lecture**

Both relations fit in main memory

- Internal memory join algorithms
- **Nested-loop join:** check for every record in R and every record in S ; time = $O(|R||S|)$
- **Sort-merge join:** sort R and S followed by merging; time = $O(|S| \log |S|)$ (if $|R| < |S|$)
- **Hash join:** build a hash table for R ; for every record in S , probe the hash table; time = $O(|S|)$ (if $|R| < |S|$)

External memory join algorithms

- At least one relation does not fit into main memory
- I/O access is the dominant cost
 - $B(R)$: number of blocks of R .
 - $|R|$ or $T(R)$: number of tuples in R .
- Memory requirement
 - M : number of blocks that fit in main memory
- **Example:** internal memory join algorithms : $B(R) + B(S)$
- We do not consider the cost of writing the output.
 - The results may be pipelined and never written to disk.

Nested-loop join of R and S

- For each block of R , and for each tuple r in the block:
 - For each block of S , and for each tuple s in the block:
 - Output rs if join condition evaluates to true over r and s
- R is called the outer table; S is called the inner table
- **cost:** $B(R) + |R| \cdot B(S)$
- **Memory requirement:** 4 (if *double buffering* is used)
- block-based nested-loop join
 - For each block of R , and for each block of S :
For each r in the R block, and for each s in the S block: ...
- **cost:** $B(R) + B(R) \cdot B(S)$
- **Memory requirement:** 4 (if *double buffering* is used)

Improving nested-loop join

- Use up the available memory buffers M
- Read $M - 2$ blocks from R
- Read blocks of S one by one and join its tuples with R tuples in main memory
- **Cost:** $B(R) + \lceil B(R) / (M - 2) \rceil B(S)$
 - almost $B(R) B(S) / M$
- **Memory requirement:** M


Index-based (zig-zag) join

- Join R and S on $R.A = S.B$
- Use ordered indexes over R.A and S.B to join the relations.
 - B+ tree
 - Use current indexes or build new ones.
 - Cost: $B(R) + B(S)$
- Memory requirement?

Index-based join algorithm


- Only R has an index over the join attribute.
- Read S, for each tuple of S find matching tuples in R.
- If S does not share its blocks with other relations
 - $V(R,A)$: Number of distinct values of attribute A in R.
 - Clustered index on R: $B(S) + T(S) B(R) / V(R,A)$.
 - Unclustered index on R: $B(S) + T(S) T(R) / V(R,A)$.
- Efficiency
 - If S is small or $V(R,A)$ is very large, not need to examine all tuples in R.
 - more efficient than nested-loop.

Two pass, multi-way merge sort



- **Problem:** sort relation R that does not fit in main memory
- Phase 1: Read R in groups of M blocks, sort, and write them as **runs** of size M on disk.

Two pass, multi-way merge sort



- Phase 2: Merge $M - 1$ blocks at a time and write the results to disk.
 - Read one block from each run.
 - Keep one block for the output.

Two pass, multi-way merge Sort

- **Cost:** $2B(R)$ in the first pass + $B(R)$ in the second pass.
- **Memory requirement:** M
 - $B(R) \leq M(M - 1)$ or simply $B(R) \leq M^2$

General multi-way merge sort

- Input: 1, 7, 4, 5, 2, 8, 9, 6, 3, 0
- Each block holds one number, and memory has 3 blocks
- Pass 0
 - 1, 7, 4 \rightarrow 1, 4, 7
 - 5, 2, 8 \rightarrow 2, 5, 8
 - 9, 6, 3 \rightarrow 3, 6, 9
 - 0 \rightarrow 0
- Pass 1
 - 1, 4, 7 + 2, 5, 8 \rightarrow 1, 2, 4, 5, 7, 8
 - 3, 6, 9 + 0 \rightarrow 0, 3, 6, 9
- Pass 2 (final)
 - 1, 2, 4, 5, 7, 8 + 0, 3, 6, 9 \rightarrow 0, 1, 2, 3, 4, 5, 6, 7, 8, 9

General multi-way merge sort

- Pass 0: read M blocks of R at a time, sort them, and write out a level-0 run
 - There are $\lceil B(R) / M \rceil$ level-0 sorted runs
- Pass i : merge $(M - 1)$ level- $(i-1)$ runs at a time, and write out a level- i run
 - $(M - 1)$ memory blocks for input, 1 to buffer output
 - # of level- i runs = # of level- $(i-1)$ runs / $(M - 1)$
- Final pass produces 1 sorted run

Analysis of general multi-way merge sort

- Number of passes: $\lceil \log_{M-1} \lceil B(R) / M \rceil \rceil + 1$
- **cost**
 - $\# \text{passes} \cdot 2 \cdot B(R)$: each pass reads the entire relation once and writes it once
 - Subtract $B(R)$ for the final pass
 - Simply $O(B(R) \cdot \log_M B(R))$
- **Memory requirement: M**


Sort-merge join algorithm

- Sort R and S according to the join attribute, then merge them
 - sort R and S using two pass multi-way merge sort
 - r, s = the first tuples in sorted R and S
 - Repeat until one of R and S is exhausted:
 - If $r.A > s.B$ then s = next tuple in S
 - else if $r.A < s.B$ then r = next tuple in R
 - else output all matching tuples, and
 - r, s = next in R and S
- **Cost:** sorting + 2 $B(R)$ + 2 $B(S)$
- What if more than M blocks match on join attribute?
 - use nested loop join algorithm
 - $B(R) B(S)$ if everything joins
- **Memory Requirement:** $B(R) \leq M^2$, $B(S) \leq M^2$

Optimized sort-merge join algorithm

- Combine join with the merge phase of sort
 - Sort R and S in M runs (overall) of size M on disk.
 - Merge and join the tuples in one pass.

Runs of R and S



Optimized sort-merge join algorithm


- **Cost:** $3B(R) + 3B(S)$
- **Memory Requirement:** $B(R) + B(S) \leq M^2$
 - because we merge them in one pass
- More efficient but more strict requirement.

(Partitioned) Hash join on R and S

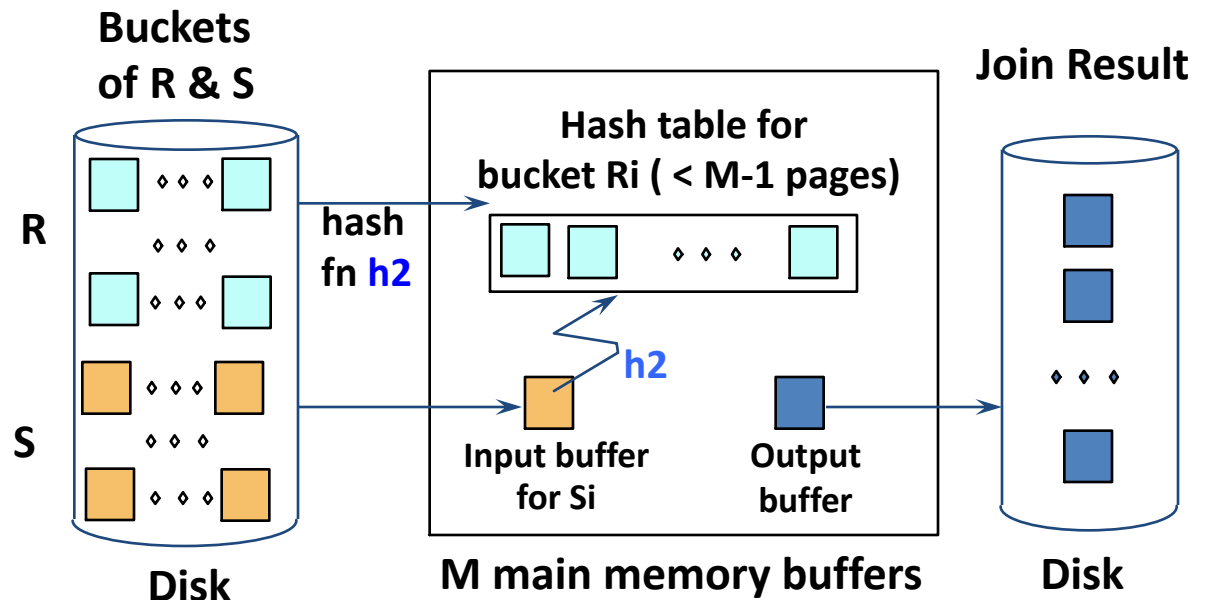
- Step 1:
 - Hash S into M buckets
 - send all buckets to disk
- Step 2
 - Hash R into M buckets
 - Send all buckets to disk
- Step 3
 - Join corresponding buckets
 - If tuples of R and S are not assigned to corresponding buckets, they do not join

Hash Join

- Partition both relations using hash fn **h**: R tuples in partition *i* will only match S tuples in partition *i*.

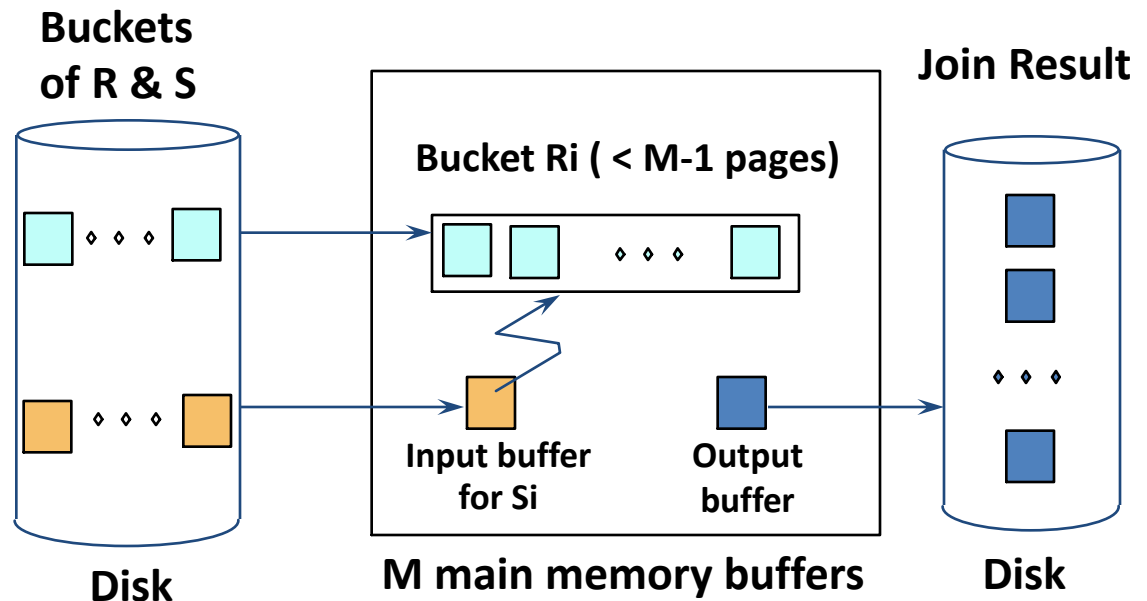


- Read in a partition of R, hash it using **h2** (<> **h**!). Scan matching partition of S, search for matches.



Hash join

- **Cost:** $3 B(R) + 3 B(S)$.
- **Memory Requirement:**
 - The smaller bucket must fit in main memory.
 - Let $\min(B(R), B(S)) = B(R)$
 - $B(R) / (M - 1) \leq M$, roughly $B(R) \leq M^2$



Handle partition overflow

- **Overflow on disk:** an R partition is larger than memory size
 - Solution: recursive partition.

Hash-based versus sort-based join

- **Hash join:** smaller amount of main memory
 - $\text{sqrt}(\min(B(R), B(S))) < \text{sqrt}(B(R) + B(S))$
 - Hash join wins if the relations have different sizes
- Hash join performance depends on the quality of hashing
 - Hard to generate balanced buckets
- Sort-based join wins if the relations are in sorted order
- Sort-based join generates sorted results
 - useful when there is **Order By** in the query
 - useful the following operators need sorted input
- Sort-based join can handle inequality join predicates

Duality of Sort and Hash

- Divide-and-conquer paradigm
- Handling very large inputs
 - Sorting: multi-level merge
 - Hashing: recursive partitioning