

Lecture 2: Query Containment

Instructor: Paris Koutris

2016

In this lecture we will study when we can say that two conjunctive queries that are syntactically different express the same query, or when the output of one query is always contained in the output of another query.

Definition 2.1 (Query Equivalence). *Two queries q_1, q_2 are equivalent, denoted $q_1 \equiv q_2$, if for every database instance I , we have $q_1(I) = q_2(I)$.*

Definition 2.2 (Query Containment). *We say that query q_1 is contained in query q_2 , denoted $q_1 \subseteq q_2$, if for every database instance I , we have $q_1(I) \subseteq q_2(I)$.*

It is straightforward to see that if $q_1 \subseteq q_2$ and $q_2 \subseteq q_1$, then $q_1 \equiv q_2$. Notice that when both queries are boolean queries, query containment becomes equivalent to logical implication!

Exercise 2.3. *Decide whether $q \subseteq q'$ for the following pairs of CQs:*

$$\begin{array}{ll}
 q_1(x, y) : -R(x, y), S(y, y), R(y, w) & q'_1(x, y) : -R(x, y), S(y, z), R(z, w) \\
 q_2(x) : -R(x, y), R(y, z), R(z, w) & q'_2(x) : -R(x, y), R(y, z) \\
 q_3(x) : -R(x, y), R(y, z), R(z, x) & q'_3(x) : -R(x, y), R(y, x) \\
 q_4(x) : -R(x, u), R(u, u) & q'_4(x) : -R(x, u), R(u, v), R(v, w)
 \end{array}$$

Is there a procedure that allows us to check whether a CQ is contained in another CQ? The answer to this question was provided by Chandra and Merlin [CM77]. But first, we need to introduce some more concepts.

Definition 2.4 (Canonical Database). *Given a conjunctive query q , the canonical database $D[q]$ is the database instance where each atom in q becomes a fact in the instance.*

As an example, the canonical database for query q_1 of the running example is the instance $D[q_1] = \{R(x, y), S(y, y), R(y, w)\}$.

Definition 2.5 (Homomorphism). *A homomorphism h from q_2 to q_1 is a function $h : \text{var}(q_2) \rightarrow \text{var}(q_1) \cup \text{const}(q_1)$ such that:*

1. *for every atom $R(x_1, x_2, \dots)$ in q_2 , there is an atom $R(h(x_1), h(x_2), \dots)$ in q_1 .*
2. *$h(\text{head}(q_2)) = \text{head}(q_1)$, where **head** denotes the head variables of a query.*

Another term used for such a homomorphism is *containment mapping*.

Example 2.6. *Consider the queries q_1, q'_1 from the running example. Consider the function h such that $h(x) = x, h(y) = y, h(z) = y, h(w) = w$, and observe that it is a homomorphism from q'_1 to q_1 .*

2.1 The Homomorphism Theorem

We can now state the central theorem for query containment.

Theorem 2.7 ([CM77]). *Given two conjunctive queries q_1, q_2 , the following statements are equivalent:*

1. $q_1 \subseteq q_2$.
2. *There exists a homomorphism h from q_2 to q_1 .*
3. $\mathbf{head}(q_1) \in q_2(D[q_1])$.

Proof. $2 \implies 1$. Let h be a homomorphism from q_2 to q_1 . Consider a database instance I . For a tuple $t \in q_1(I)$, we need to prove that $t \in q_2(I)$. Since $t \in q_1(I)$, there exists a valuation v such that $t = v(\mathbf{head}(q_1))$. Consider the composition of v with the homomorphism h : $g = v \circ h$. We will show that g is a valuation for query q_2 . Consider now an atom $R(y_1, \dots, y_m)$ in q_2 . By the definition of homomorphism, $R(h(y_1), \dots, h(y_m))$ is in q_1 , and then by definition of the valuation, $R(g(y_1), \dots, g(y_m))$ is a fact of I . Also, $g(\mathbf{head}(q_2)) = v(h(\mathbf{head}(q_2))) = v(\mathbf{head}(q_1)) = t$.

$1 \implies 3$. Let $q_1 \subseteq q_2$, and consider the canonical database $D[q_1]$. Then, $\mathbf{head}(q_1) \in q_1(D[q_1])$, since the identity function is a valuation. Since $q_1 \subseteq q_2$, $\mathbf{head}(q_1) \in q_2(D[q_1])$ as well.

$3 \implies 2$. Since $\mathbf{head}(q_1) \in q_2(D[q_1])$, there exists a valuation v for query q_2 . This valuation is a homomorphism from q_2 to q_1 , since (a) $v(\mathbf{head}(q_1)) = \mathbf{head}(q_2)$ and (b) for every atom $R(y_1, \dots, y_m)$, $R(v(y_1), \dots, v(y_m))$ is a fact in the canonical database $D[q_1]$, and so an atom in q_1 . \square

Exercise 2.8. Let $q_1(x) : -R(x, y), R(y, z), R(z, w)$ and $q_2(x) : -R(x, y), R(y, x)$. Is $q_2 \subseteq q_1$?

Theorem 2.9 ([CM77]). *The problem of query containment for conjunctive queries is NP-complete.*

Proof. The NP-hardness is proved by reduction from graph 3-colorability: given a graph $G(V, E)$, is it possible to color the vertices with 3 colors, such that every pair of neighboring vertices has different colors? Given the graph G , consider the binary relation $R(x, y)$ that represents the edges of the graph, and let q_G be the boolean conjunctive query that corresponds to the graph. For example, if the graph has the edges $(a, b), (b, c), (c, d), (c, a)$, then the query is $q_G() : -R(a, b), R(b, c), R(c, d), R(c, a)$. Consider now the query $K_3() : -R(x, y), R(y, z), R(z, x)$. It is easy to see that G is 3-colorable if and only if there exists a homomorphism from q_G to K_3 . By the homomorphism theorem, G is 3-colorable if and only if $K_3 \subseteq q_G$.

The membership in NP follows again from the homomorphism theorem. To decide whether $q_1 \subseteq q_2$, we need to find a homomorphism from q_2 to q_1 . If we guess a function, we can check whether it is a homomorphism in polynomial time. \square

As a corollary of the homomorphism theorem, we can also pinpoint the complexity of evaluating a conjunctive query.

Corollary 2.10. *The problem of evaluating a conjunctive query is NP-complete.*

It can be shown that query equivalence for conjunctive queries is also *NP*-complete. Even though query containment is an *NP*-complete problem, it is feasible problem to solve because typically the size of the query is not very large.

2.2 Query Minimization

Here we discuss how we can minimize a given conjunctive query. In other words, given a conjunctive query q , can we find an equivalent CQ q' such that it has as few atoms as possible?

Definition 2.11 (Minimal Query). *A conjunctive query q is minimal if for every other conjunctive query q' such that $q \equiv q'$, q' has at least as many atoms as q .*

Exercise 2.12. *Is the query $q(x) : -R(x,y), R(x,z), R(z,w)$ minimal? What about the query $q(x) : -R(x,a), R(x,y), R(y,z)$?*

Find the minimal equivalent CQ to $q(x,y) : -R(y,x), R(z,x), R(w,x), R(x,u)$.

Using query minimization to evaluate CQs sounds like a promising approach. Instead of evaluating a CQ q directly, we can first compute a minimal equivalent CQ q' , and then evaluate the new query q' to obtain the desired answer faster.

Theorem 2.13. *Let q be a conjunctive query.*

1. *There exists a minimal equivalent conjunctive query q' that can be obtained from q by removing zero or more atoms.*
2. *All minimal equivalent queries of q are isomorphic.*

Proof. (1) Let q'' be a minimal equivalent query to q . Then, by the homomorphism theorem, there exists a homomorphism h from q to q'' , and g from q'' to q . Let q' be the query that results if we apply $h \circ g$ to q . It is straightforward to verify that $q' \equiv q$ and that q' has at most as many atoms as q'' .

(2) This item tells us that, even though there is not a unique minimal equivalent query, minimal queries are the same up to renaming of variables. The proof is left as an exercise! \square

Below is a simple procedure to compute a minimal equivalent query:

- Choose an atom from q and remove it to obtain a new query q' . We know from the homomorphism theorem that $q \subseteq q'$.
- Check if $q' \subseteq q$; if so, then q' is equivalent and we can continue the process of removing another atom.
- If not, try to remove another atom from q .

Unfortunately, CQ minimization is an *NP*-hard problem, so we cannot hope to have a general efficient algorithm. Even though CQ minimization looks like a promising avenue, in practice query optimizers do not use it (why?).

2.3 Beyond Conjunctive Queries

Query containment for all of relational algebra (so the first-order fragment of relational calculus) is actually an undecidable problem. To prove undecidability, we can reduce from the problem of finite satisfiability of first-order formulas. A first order sentence ϕ is *finitely satisfiable* if there exists a finite database instance I such that ϕ is true over I .

Theorem 2.14 (Trakhtenbrot's Theorem [T50]). *Finite satisfiability is undecidable in first-order logic.*

To show the reduction, consider a first-order sentence ϕ and construct the following two relational calculus queries: $q_1 = \exists x(R(x) \wedge \phi)$ and $q_2 = \exists x(R(x) \wedge (x \neq x))$. It is clear that q_2 is always false, hence $q_1 \subseteq q_2$ if and only if q_1 is always false, which is equivalent to ϕ being not finitely satisfiable.

However, as we will see next, query containment is more tractable for classes of queries that are between CQs and the full relational calculus. Let's consider first the class of UCQs.

Theorem 2.15 ([SY80]). *Let $q = q_1 \cup q_2 \cup \dots \cup q_m$ and $q' = q'_1 \cup q'_2 \cup \dots \cup q'_n$ be UCQs. The following statements are equivalent:*

1. $q \subseteq q'$.
2. For every $i = 1, \dots, m$, there exists some $j = 1, \dots, n$ such that $q_i \subseteq q'_j$.

Proof. $2 \implies 1$. Straightforward. Consider an instance I and let $t \in q(I)$. Then, $t \in q_i(I)$ for some $i = 1, \dots, m$, and since $q_i \subseteq q'_j$ for some j , $t \in q'_j(I)$. Thus, $t \in q'(I)$.

$1 \implies 2$. Consider the canonical instance $D[q_i]$. We then have: $q_i(D[q_i]) \subseteq q(D[q_i]) \subseteq q'(D[q_i])$. It must be then that $\text{head}(q_i) \in q'_j(D[q_i])$, which by the homomorphism theorem implies that $q_i \subseteq q'_j$. \square

As a corollary of the above result, we obtain that the problem of query containment for unions of conjunctive queries is also an *NP*-complete problem.

The complexity for query containment for $CQ^\neq, CQ^<$ is much higher than *NP*-complete: it is in the class Π_2^P . We won't get into details for this complexity result.

References

[Alice] S. ABITEBOUL, R. HULL and V. VIANU, "Foundations of Databases."

- [CM77] A.K. CHANDRA and P.M. MERLIN, "Optimal implementation of conjunctive queries in relational data bases," *STOC 1977*.
- [SY80] Y. SAGIV and M. YANNAKAKIS, "Equivalences Among Relational Expressions with the Union and Difference Operators", *JACM 1980*.
- [T50] B.A. TRAKHTENBROT, "The impossibility of an algorithm for the decidability problem on finite classes", *Doklady AN SSR 70(4)*, 569?572, 1950.