

ONID:

CS540 Database Management Systems
Winter 2021
School of Electrical Engineering & Computer Science
Oregon State University
Midterm Examination
Time Limit: 120 minutes

- Print your name and ONID below. In addition, print your ONID in the upper right corner of every page.

Name and ONID: _____

- The exam is open book and notes.
- Any form of cheating on the examination will result in a zero grade.
- The submissions should be in pdf format.
- Please make your answers clear and succinct; you will lose credit for verbose, convoluted, or confusing answers. *Simplicity does count!*

Question:	1	2	3	4	Total
Points:	4	5	6	5	20
Score:					




Figure 1: The B+tree for question 1.

1. B+tree Index

Consider the B+tree with degree, i.e., order, $d = 1$ shown in Figure 1. The picture shows both keys and pointers in each node. We do *not* show the links to the data blocks on the underlying relation from the leaf nodes of the B+tree.

- (a) (4 points) Show the B+tree produced after inserting search keys with values 25 and 30. You should first insert 25 and then 30.

2. Indexing

- (a) (5 points) Consider the relation `Employee(id, name, salary)` and the following SQL query.

```
SELECT name
FROM Employee
WHERE salary > 1000 and salary < 10000
```

What index will improve the running time of this query more than any other index? You must explain the attribute(s) of the index, its type (B+tree or hash), and whether it is clustered. You should justify your answer to each part.

3. Query Processing: Join Algorithms

Consider relations $Student(ID, Name)$ and $Enrollment(CourseID, SID)$. The size of relation $Student$ is 40,000 blocks and the size of $Enrollment$ is 400 blocks. Assume that the query processor has to choose between the internal memory (in-memory) join algorithms, block-based nested loop algorithm, and the sort-merge join with two-pass multi-way merge-sort algorithm for sorting to perform the join of $Student \bowtie_{ID=SID} Enrollment$.

- (a) (6 points) Given that we have only 4 blocks available in main memory, i.e., $M = 4$, which one of the aforementioned algorithms is the fastest one to perform the join? You should also provide the cost, i.e., number of I/O accesses, of the join using your proposed algorithm.

4. Query Processing

The *left anti-join* operator is a variant of the relational join. The basic idea is to output only those tuples in the left relation that do *not* join with any tuple in the right relation. For example, consider two relations $R(A, B)$ and $S(B, C)$. Assume that R has the following tuples: $(1, 10)$, $(1, 20)$, $(2, 30)$, $(2, 40)$ and S has the following tuples: $(10, 75)$, $(10, 85)$, $(30, 95)$. The left anti-join of R and S on attribute B will produce the following tuples: $(1, 20)$ and $(2, 40)$.

Given the above description of left anti-joins, answer the following question.

- (a) (5 points) How can you adapt the sort-merge join algorithm to process left anti-joins? Describe only essential modifications to the algorithm. You may use the two-pass multi-way merge-sort algorithm for sorting.